

Module 5: Kokkos Kernels Math Library

August 27, 2020

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.
SAND2020-7475 PE

Online Resources:

- ▶ <https://github.com/kokkos>:
 - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
 - ▶ Slides, recording and Q&A for the Lectures
- ▶ <https://github.com/kokkos/kokkos/wiki>:
 - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
 - ▶ Slack channel for Kokkos.
 - ▶ Please join: fastest way to get your questions answered.
 - ▶ Can whitelist domains, or invite individual people.

- ▶ 07/17 Module 1: Introduction, Building and Parallel Dispatch
- ▶ 07/24 Module 2: Views and Spaces
- ▶ 07/31 Module 3: Data Structures + MultiDimensional Loops
- ▶ 08/07 Module 4: Hierarchical Parallelism
- ▶ 08/14 Module 5: Tasking, Streams and SIMD
- ▶ 08/21 Module 6: Internode: MPI and PGAS
- ▶ 08/28 Module 7: Tools: Profiling, Tuning and Debugging
- ▶ **09/04 Module 8: Kernels: Sparse and Dense Linear Algebra**
- ▶ 09/11 Reserve Day

Tools Stuff

Dense Linear Algebra (BLAS)

- ▶ BLAH.
- ▶ More BLAH.

Sparse Linear Algebra

- ▶ Sparse BLAH.
- ▶ Sparse BLAH2.

Graph Functions

- ▶ Graph BLAH.

Dense Linear Algebra (BLAS)

Kokkos Kernels dense linear algebra capabilities.

Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

How do we make Kokkos and Fortran talk with each other?

Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
 - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
 - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


How do we create an nd_array_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



Sparse Linear Algebra

Sparse linear algebra data structures and functions.

Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

How do we make Kokkos and Fortran talk with each other?

Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
 - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
 - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


How do we create an nd_array_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



Sparse Solvers

Sparse linear algebra data structures and functions.

Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

How do we make Kokkos and Fortran talk with each other?

Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
 - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
 - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


How do we create an nd_array_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



Graph Kernels

Kokkos Kernels functionality for graph computations.

Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
 - ▶ Keep our e.g. Fortran mains, drivers, physics packages
 - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

How do we make Kokkos and Fortran talk with each other?

Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
 - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
 - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

nd_array_t

The compatibility glue between Fortran arrays and Kokkos Views.

Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


How do we create an nd_array_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

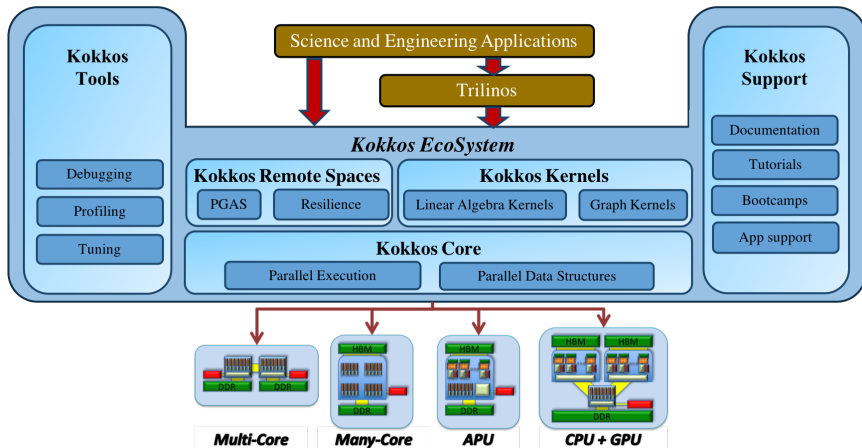

3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



Summary

More Summary





Kokkos Core:	C.R.Trott , J. Ciesko, V. Dang, N. Ellingwood, D.S. Hollman, D. Ibanez, J. Miles, J. Wilke, , H. Finkel, N. Liber, D. Lebrun-Grandie, D. Arndt, B. Turcksin, J. Madsen, R. Gayatri former: H.C. Edwards, D. Labreche, G. Mackey, S. Bova, D. Sunderland
Kokkos Kernels:	S. Rajamanickam , L. Berger, V. Dang, N. Ellingwood, E. Harvey, B. Kelley, K. Kim, C.R. Trott, J. Wilke, S. Acer
Kokkos Tools	D. Poliakoff , C. Lewis, S. Hammond, D. Ibanez, J. Madsen, S. Moore, C.R. Trott
Kokkos Support	C.R. Trott , G. Shipmann, G. Womeldorff, and all of the above former: H.C. Edwards, G. Lopez, F. Foertter

Online Resources:

- ▶ <https://github.com/kokkos>:
 - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
 - ▶ Slides, recording and Q&A for the Lectures
- ▶ <https://github.com/kokkos/kokkos/wiki>:
 - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
 - ▶ Slack channel for Kokkos.
 - ▶ Please join: fastest way to get your questions answered.
 - ▶ Can whitelist domains, or invite individual people.