

# The Kokkos Lectures

## Module 5: Kokkos Kernels Math Library

August 27, 2020

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.  
SAND2020-7475 PE

## Online Resources:

- ▶ <https://github.com/kokkos>:
  - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
  - ▶ Slides, recording and Q&A for the Lectures
- ▶ <https://github.com/kokkos/kokkos/wiki>:
  - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
  - ▶ Slack channel for Kokkos.
  - ▶ Please join: fastest way to get your questions answered.
  - ▶ Can whitelist domains, or invite individual people.

- ▶ 07/17 Module 1: Introduction, Building and Parallel Dispatch
- ▶ 07/24 Module 2: Views and Spaces
- ▶ 07/31 Module 3: Data Structures + MultiDimensional Loops
- ▶ 08/07 Module 4: Hierarchical Parallelism
- ▶ 08/14 Module 5: Tasking, Streams and SIMD
- ▶ 08/21 Module 6: Internode: MPI and PGAS
- ▶ 08/28 Module 7: Tools: Profiling, Tuning and Debugging
- ▶ **09/04 Module 8: Kernels: Sparse and Dense Linear Algebra**
- ▶ 09/11 Reserve Day

## Tools Stuff

## **Dense Linear Algebra (BLAS)**

- ▶ BLAH.
- ▶ More BLAH.

## **Sparse Linear Algebra**

- ▶ Sparse BLAH.
- ▶ Sparse BLAH2.

## **Graph Functions**

- ▶ Graph BLAH.

# Dense Linear Algebra (BLAS)

Kokkos Kernels dense linear algebra capabilities.

## **Learning objectives:**

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

## HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **HPC world owns many Fortran LOC!**

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **How do we make Kokkos and Fortran talk with each other?**



## Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
  - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
  - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

### FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


### How do we create an nd\_array\_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

### 3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



# Sparse Linear Algebra

Sparse linear algebra data structures and functions.

## Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

## HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **HPC world owns many Fortran LOC!**

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **How do we make Kokkos and Fortran talk with each other?**

## Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
  - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
  - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

### FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.



## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


### How do we create an nd\_array\_t?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

### 3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



# Sparse Solvers

Sparse linear algebra data structures and functions.

## Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

# Sparse Solvers + Preconditioning

- ▶ Basic idea is to use the Krylov subspace method on a modified system to increase the convergence rate
- ▶ Common Krylov subspace solvers: Conjugate Gradient (CG), Biconjugate Gradient Stabilized (BiCGSTAB), Generalized Minimal Residual (GMRES)
- ▶ Common preconditioners: Jacobi (or diagonal), Sparse Approximate Inverse (SPAI), Incomplete LU factorization (ILU)

## HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **HPC world owns many Fortran LOC!**

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

**How do we make Kokkos and Fortran talk with each other?**

## Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
  - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
  - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

### FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.



## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### **Keeps Track of:**

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation

## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


### How do we create an `nd_array_t`?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

### 3D Heat Conduction

- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



# Graph Kernels

Kokkos Kernels functionality for graph computations.

## Learning objectives:

- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views
- ▶ Calling BLAS functions with Views

## HPC world owns many Fortran LOC!

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

## **HPC world owns many Fortran LOC!**

- ▶ We generally cannot port it all at once.
- ▶ We need an incremental porting strategy
  - ▶ Keep our e.g. Fortran mains, drivers, physics packages
  - ▶ But port relevant infrastructure, or hotspot kernels to C++ and Kokkos

**How do we make Kokkos and Fortran talk with each other?**

## Fortran Language Compatibility Layer (FLCL)

- ▶ Pass multidimensional arrays accross the C++/Fortran boundary
  - ▶ See Fortran arrays as Kokkos Views and vice versa
- ▶ Create Kokkos View and DualView from Fortran
  - ▶ Allows Fortran to be the memory owner but call C++ functions with Kokkos kernels for CUDA/HIP
- ▶ Initialize and Finalize Kokkos from Fortran
- ▶ FortranIndex<T> scalar type to deal with 1 vs 0 based indexing in sparse data structures

### FLCL

The Fortran Language Compatibility Layer allows an incremental porting of a Fortran code to Kokkos.

## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### **Keeps Track of:**

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation



## nd\_array\_t

The compatibility glue between Fortran arrays and Kokkos Views.

### Keeps Track of:

- ▶ An array's rank
- ▶ Extents of the array
- ▶ Strides of the array
- ▶ Pointer to the allocation


### How do we create an `nd_array_t`?

- ▶ Explicit routines like `to_nd_array_i64_6`
- ▶ Simple interface taking a fortran array as argument

```
array = to_nd_array(foo);
```

### 3D Heat Conduction

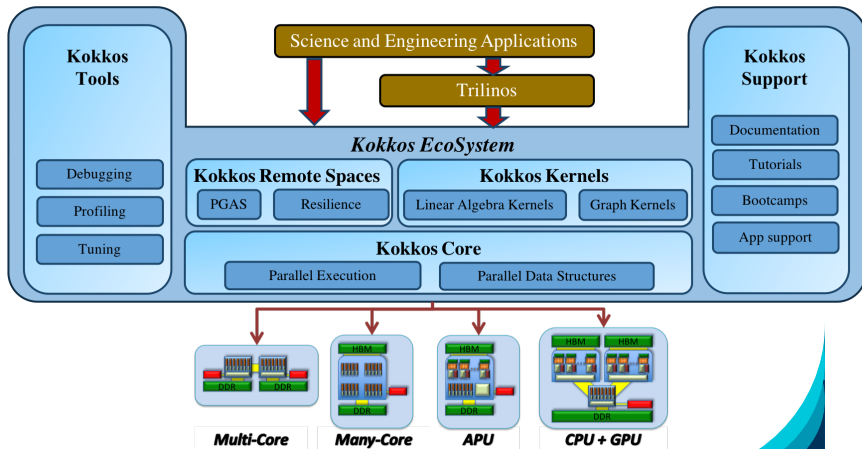
- ▶ Heat conduction inside the body
- ▶ Thermal radiation (Black Body) on surface
- ▶ Incoming power flow from one direction



## Summary

## More Summary







|                        |   |
|------------------------|---|
| <b>Kokkos Core:</b>    | <b>C.R.Trott</b> , J. Ciesko, V. Dang, N. Ellingwood, D.S. Hollman, D. Ibanez, J. Miles, J. Wilke, , H. Finkel, N. Liber, D. Lebrun-Grandie, D. Arndt, B. Turcksin, J. Madsen, R. Gayatri<br>former: H.C. Edwards, D. Labreche, G. Mackey, S. Bova, D. Sunderland |
| <b>Kokkos Kernels:</b> | <b>S. Rajamanickam</b> , L. Berger, V. Dang, N. Ellingwood, E. Harvey, B. Kelley, K. Kim, C.R. Trott, J. Wilke, S. Acer   |
| <b>Kokkos Tools</b>    | <b>D. Poliakoff</b> , C. Lewis, S. Hammond, D. Ibanez, J. Madsen, S. Moore, C.R. Trott  |
| <b>Kokkos Support</b>  | <b>C.R. Trott</b> , G. Shipmann, G. Womeldorff, and all of the above<br>former: H.C. Edwards, G. Lopez, F. Foertter   |

## Online Resources:

- ▶ <https://github.com/kokkos>:
  - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
  - ▶ Slides, recording and Q&A for the Lectures
- ▶ <https://github.com/kokkos/kokkos/wiki>:
  - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
  - ▶ Slack channel for Kokkos.
  - ▶ Please join: fastest way to get your questions answered.
  - ▶ Can whitelist domains, or invite individual people.